

Deep Generative Models: Part I

Q. Bertrand^{1, 3}, M. Massias^{1, 2}, R. Emonet^{3, 1}

¹Inria Lyon

²ENS Lyon

³Laboratoire Hubert Curien - Saint-Étienne

April 11, 2025

A Few Words About Me

Academic

- ▶ 2018 - 2021 Ph. D. at Inria Saclay: Large-Scale Optimization for Neuroimaging
- ▶ 2021 - 2024 Postdoc at Mila and Université de Montréal
- ▶ 2024 - now Researcher at Inria Lyon, based in Laboratoire Hubert Curien

A Few Words About Me

Academic

- ▶ 2018 - 2021 Ph. D. at Inria Saclay: Large-Scale Optimization for Neuroimaging
- ▶ 2021 - 2024 Postdoc at Mila and Université de Montréal
- ▶ 2024 - now Researcher at Inria Lyon, based in Laboratoire Hubert Curien

Research interests:

- ▶ Numerical optimization
- ▶ Generative optimization
 - Retraining of generative models on their own data
- ▶ Multi-agent reinforcement learning
 - Emergence of cooperation in multi-agent settings

Outline

What are Generative Models?

Generative Adversarial Networks

Normalizing Flows

Continuous Normalizing Flow

Conditional Flow Matching & Diffusion

Some of our own Research (if Time?)

References

What are Generative Models?

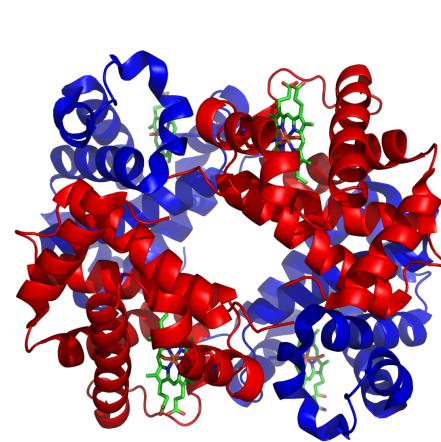
Generative Models are Powerful



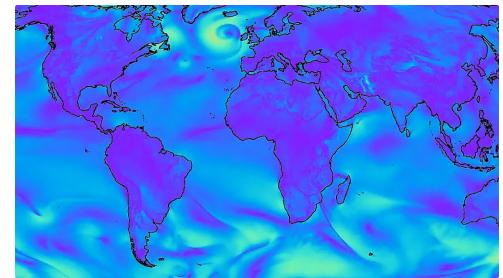
Text-to-Image



Text-to-Text



Protein Design



Weather Forecast

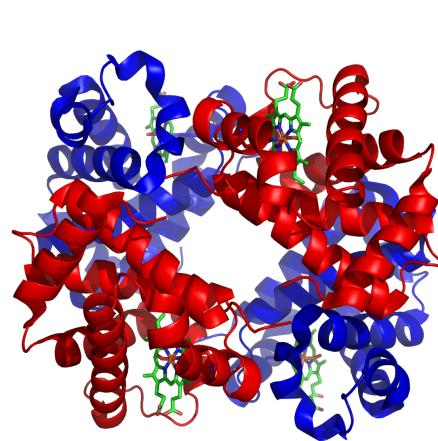
Generative Models are Powerful



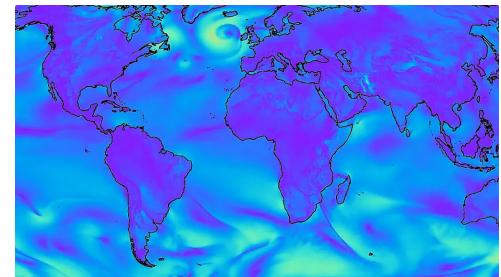
Text-to-Image



Text-to-Text



Protein Design



Weather Forecast

Academic References:

- ▶ [Stability AI \(2023\)](#)
- ▶ [Watson et al. \(2023\)](#)
- ▶ [Lam et al. \(2023\)](#)
- ▶ [Xu et al. \(2025\)](#)

What are Generative Models 1/3

Data: $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}$



Goal: new synthetic samples $\tilde{\mathbf{x}}^{(1)}, \dots, \tilde{\mathbf{x}}^{(n)}$



Generative Models 101:

- ▶ **Setting:** Access to samples $(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}) \rightarrow$ unlabeled data

What are Generative Models 1/3

Data: $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}$



Goal: new synthetic samples $\tilde{\mathbf{x}}^{(1)}, \dots, \tilde{\mathbf{x}}^{(n)}$



Generative Models 101:

- ▶ **Setting:** Access to samples $(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}) \rightarrow$ unlabeled data
 - Draw from a probability distribution $p, \mathbf{x}^{(i)} \sim p$

What are Generative Models 1/3

Data: $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}$



Goal: new synthetic samples $\tilde{\mathbf{x}}^{(1)}, \dots, \tilde{\mathbf{x}}^{(n)}$



Generative Models 101:

- ▶ **Setting:** Access to samples $(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}) \rightarrow$ unlabeled data
 - Draw from a probability distribution $p, \mathbf{x}^{(i)} \sim p$
 - e.g., a set of images

What are Generative Models 1/3

Data: $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}$



Goal: new synthetic samples $\tilde{\mathbf{x}}^{(1)}, \dots, \tilde{\mathbf{x}}^{(n)}$



Generative Models 101:

- ▶ **Setting:** Access to samples $(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}) \rightarrow$ unlabeled data
 - Draw from a probability distribution $p, \mathbf{x}^{(i)} \sim p$
 - e.g., a set of images
- ▶ **Goal:** Create new synthetic samples $\tilde{\mathbf{x}}^{(i)} \sim p$

What are Generative Models 1/3

Data: $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}$



Goal: new synthetic samples $\tilde{\mathbf{x}}^{(1)}, \dots, \tilde{\mathbf{x}}^{(n)}$



Generative Models 101:

- ▶ **Setting:** Access to samples $(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}) \rightarrow$ unlabeled data
 - Draw from a probability distribution $p, \mathbf{x}^{(i)} \sim p$
 - e.g., a set of images
- ▶ **Goal:** Create new synthetic samples $\tilde{\mathbf{x}}^{(i)} \sim p$
 - e.g., generate new images

What are Generative Models 2/3

Data: $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)})$, ..., $(\mathbf{x}^{(n)}, \mathbf{y}^{(n)})$



Goal: new synthetic samples $\tilde{\mathbf{x}}^{(i)} \mid \mathbf{y}^{(i)}$



Generative Models 201: Class Conditional Generative Models

- **Setting:** Access to samples $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)})$, ..., $(\mathbf{x}^{(n)}, \mathbf{y}^{(n)})$) \rightarrow labeled data

What are Generative Models 2/3

Data: $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)})$, ..., $(\mathbf{x}^{(n)}, \mathbf{y}^{(n)})$



Goal: new synthetic samples $\tilde{\mathbf{x}}^{(i)} \mid \mathbf{y}^{(i)}$



Generative Models 201: Class Conditional Generative Models

- **Setting:** Access to samples $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)})$, ..., $(\mathbf{x}^{(n)}, \mathbf{y}^{(n)})$) → labeled data
 - Draw from a probability distribution p , $\mathbf{x}^{(i)} \sim p(\cdot \mid \mathbf{y}^{(i)})$

What are Generative Models 2/3

Data: $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)})$, ..., $(\mathbf{x}^{(n)}, \mathbf{y}^{(n)})$



Goal: new synthetic samples $\tilde{\mathbf{x}}^{(i)} \mid \mathbf{y}^{(i)}$



Generative Models 201: Class Conditional Generative Models

- ▶ **Setting:** Access to samples $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)})$, ..., $(\mathbf{x}^{(n)}, \mathbf{y}^{(n)})$) → labeled data
 - Draw from a probability distribution p , $\mathbf{x}^{(i)} \sim p(\cdot \mid \mathbf{y}^{(i)})$
 - e.g., a set of images $\mathbf{x}^{(i)}$ with class $\mathbf{y}^{(i)}$

What are Generative Models 2/3

Data: $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)})$, ..., $(\mathbf{x}^{(n)}, \mathbf{y}^{(n)})$



Goal: new synthetic samples $\tilde{\mathbf{x}}^{(i)} \mid \mathbf{y}^{(i)}$



Generative Models 201: Class Conditional Generative Models

- ▶ **Setting:** Access to samples $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)})$, ..., $(\mathbf{x}^{(n)}, \mathbf{y}^{(n)})$) → labeled data
 - Draw from a probability distribution p , $\mathbf{x}^{(i)} \sim p(\cdot \mid \mathbf{y}^{(i)})$
 - e.g., a set of images $\mathbf{x}^{(i)}$ with class $\mathbf{y}^{(i)}$
- ▶ **Goal:** Create new synthetic samples $\tilde{\mathbf{x}}^{(i)} \sim p(\cdot \mid \mathbf{y}^{(i)})$

What are Generative Models 2/3

Data: $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)})$, ..., $(\mathbf{x}^{(n)}, \mathbf{y}^{(n)})$



Goal: new synthetic samples $\tilde{\mathbf{x}}^{(i)} \mid \mathbf{y}^{(i)}$



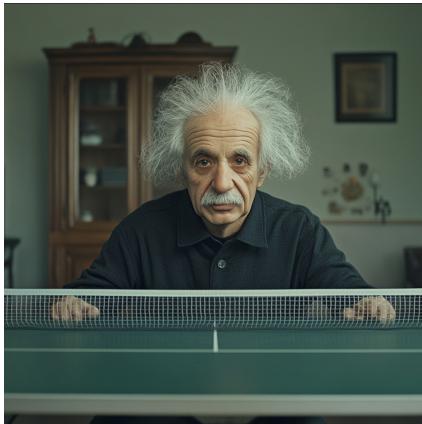
Generative Models 201: Class Conditional Generative Models

- ▶ **Setting:** Access to samples $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)})$, ..., $(\mathbf{x}^{(n)}, \mathbf{y}^{(n)})$) → labeled data
 - Draw from a probability distribution p , $\mathbf{x}^{(i)} \sim p(\cdot \mid \mathbf{y}^{(i)})$
 - e.g., a set of images $\mathbf{x}^{(i)}$ with class $\mathbf{y}^{(i)}$
- ▶ **Goal:** Create new synthetic samples $\tilde{\mathbf{x}}^{(i)} \sim p(\cdot \mid \mathbf{y}^{(i)})$
 - e.g., generate specifically new images of planes

What are Generative Models 3/3



An avocado chair



A photo of Einstein



A house made of sushi

Generative Models 301: Text-to-Image Models

- **Setting:** Access to samples $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}), \dots, (\mathbf{x}^{(n)}, \mathbf{y}^{(n)})$) → pairs (image, caption)

What are Generative Models 3/3



An avocado chair



A photo of Einstein



A house made of sushi

Generative Models 301: Text-to-Image Models

- ▶ **Setting:** Access to samples $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}), \dots, (\mathbf{x}^{(n)}, \mathbf{y}^{(n)})$) → pairs (image, caption)
 - Draw from a probability distribution p , $\mathbf{x}^{(i)} \sim p(\cdot \mid \text{caption } \mathbf{y}^{(i)})$

What are Generative Models 3/3



An avocado chair



A photo of Einstein



A house made of sushi

Generative Models 301: Text-to-Image Models

- ▶ **Setting:** Access to samples $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}), \dots, (\mathbf{x}^{(n)}, \mathbf{y}^{(n)})$) → pairs (image, caption)
 - Draw from a probability distribution p , $\mathbf{x}^{(i)} \sim p(\cdot \mid \text{caption } \mathbf{y}^{(i)})$
 - e.g., a set of images $\mathbf{x}^{(i)}$ with caption $\mathbf{y}^{(i)}$

What are Generative Models 3/3



An avocado chair



A photo of Einstein



A house made of sushi

Generative Models 301: Text-to-Image Models

- ▶ **Setting:** Access to samples $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}), \dots, (\mathbf{x}^{(n)}, \mathbf{y}^{(n)})$) → pairs (image, caption)
 - Draw from a probability distribution p , $\mathbf{x}^{(i)} \sim p(\cdot \mid \text{caption } \mathbf{y}^{(i)})$
 - e.g., a set of images $\mathbf{x}^{(i)}$ with caption $\mathbf{y}^{(i)}$
- ▶ **Goal:** Create new synthetic samples $\tilde{\mathbf{x}}^{(i)} \sim p(\cdot \mid \text{caption } \mathbf{y}^{(i)})$

What are Generative Models 3/3



An avocado chair



A photo of Einstein



A house made of sushi

Generative Models 301: Text-to-Image Models

- **Setting:** Access to samples $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}), \dots, (\mathbf{x}^{(n)}, \mathbf{y}^{(n)})$) → pairs (image, caption)
 - Draw from a probability distribution p , $\mathbf{x}^{(i)} \sim p(\cdot \mid \text{caption } \mathbf{y}^{(i)})$
 - e.g., a set of images $\mathbf{x}^{(i)}$ with caption $\mathbf{y}^{(i)}$
- **Goal:** Create new synthetic samples $\tilde{\mathbf{x}}^{(i)} \sim p(\cdot \mid \text{caption } \mathbf{y}^{(i)})$
- **Trick:** Use a “meaningful” text/caption representation

A Few Comments

A Few Comments

- ▶ **Text-to-Text Models**
 - Conditional text generation w.r.t. the previous text sequence

A Few Comments

- ▶ **Text-to-Text Models**
 - Conditional text generation w.r.t. the previous text sequence
- ▶ **Evaluation?**

A Few Comments

- ▶ **Text-to-Text Models**
 - Conditional text generation w.r.t. the previous text sequence
- ▶ **Evaluation?**
 - Quality of the generated images?

A Few Comments

- ▶ **Text-to-Text Models**
 - Conditional text generation w.r.t. the previous text sequence
- ▶ **Evaluation?**
 - Quality of the generated images?
 - No cross-validation

A Few Comments

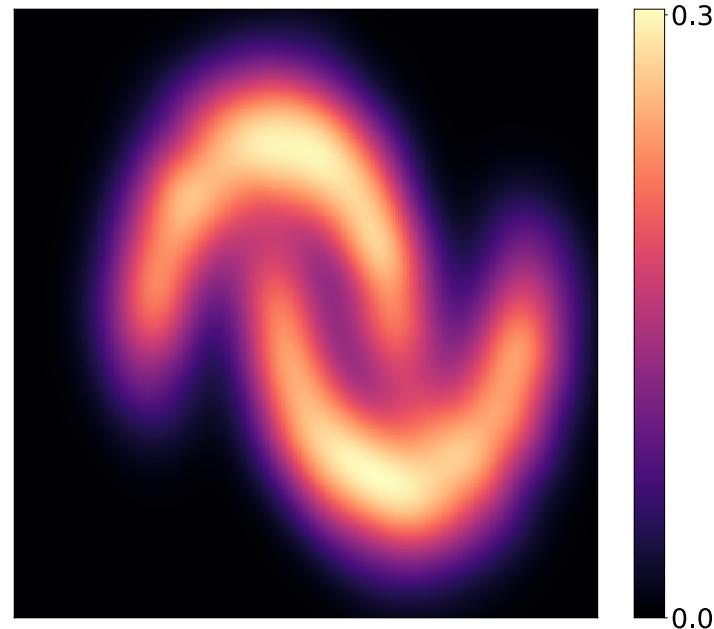
- ▶ **Text-to-Text Models**
 - Conditional text generation w.r.t. the previous text sequence
- ▶ **Evaluation?**
 - Quality of the generated images?
 - No cross-validation
 - **Very challenging** for text-to-image and text-to-text models

A Few Comments

- ▶ **Text-to-Text Models**
 - Conditional text generation w.r.t. the previous text sequence
- ▶ **Evaluation?**
 - Quality of the generated images?
 - No cross-validation
 - **Very challenging** for text-to-image and text-to-text models
 - Often use other models to assess the quality e.g., GPT

“Implicit” Generative Models: Problem Setting 1/4

Goal: estimate the unknown complex density p_{data}



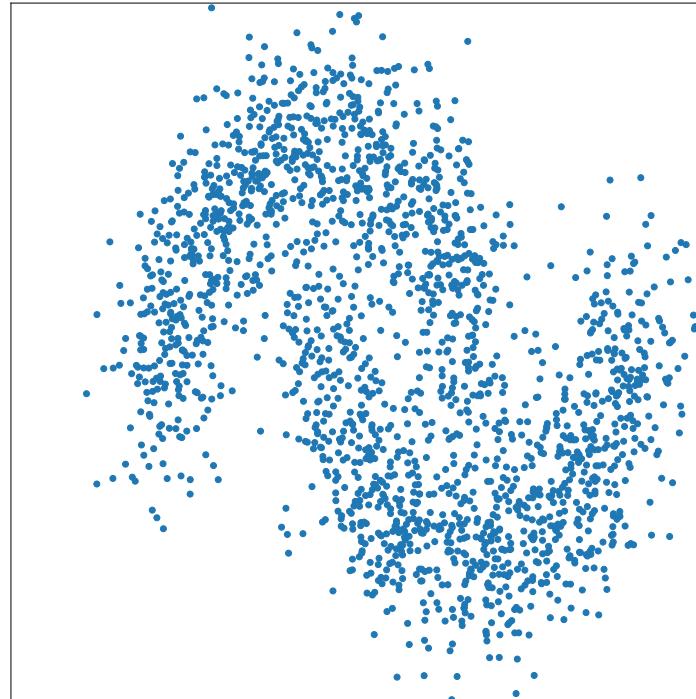
Unknown density p_{data}

“Implicit” Generative Models: Problem Setting 2/4

- ▶ **Unknown complex** data distribution p_{data}
- ▶ **Data:** what do you have access to?

“Implicit” Generative Models: Problem Setting 2/4

- ▶ **Unknown complex** data distribution p_{data}
- ▶ **Data**: what do you have access to?

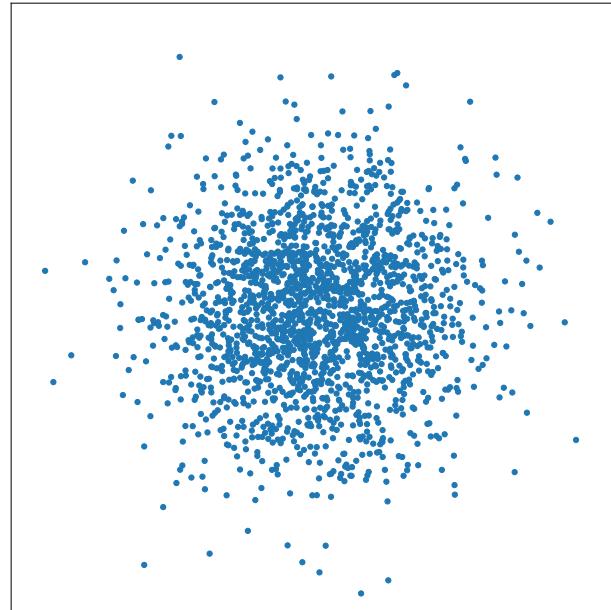


Data points $x^{(i)} \in \mathbb{R}^2, x^{(i)} \sim p_{\text{data}}$

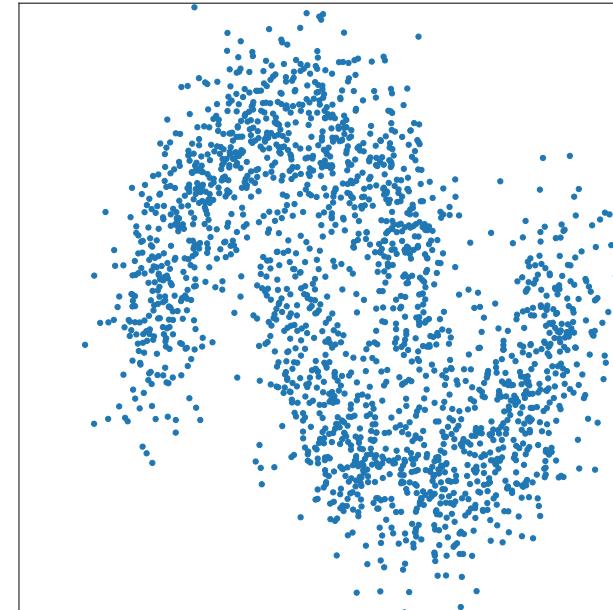
“Implicit” Generative Models: Problem Setting 3/4

- ▶ **Unknown complex** data distribution p_{data}
- ▶ **Goal:** transform a **simple** data distribution p_0 in the **complex** p_{data}

Points $x_0 \sim p_0$ drawn from a **simple** data distribution $\mathcal{N}(0, 1)$



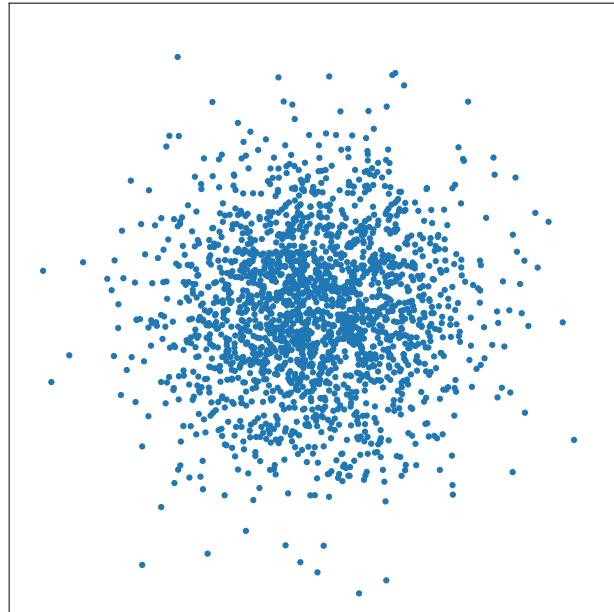
Data points $x^{(i)} \in \mathbb{R}^2$ from a **complex data distribution** $x^{(i)} \sim p_{\text{data}}$



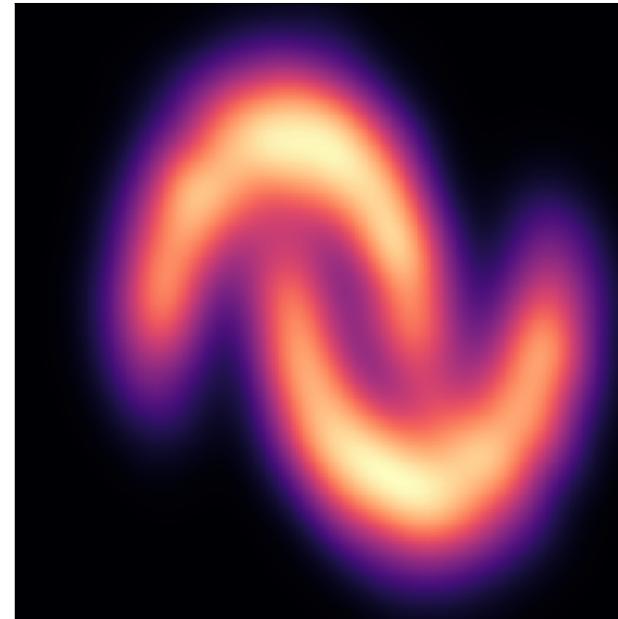
“Implicit” Generative Models: Problem Setting 4/4

- ▶ **Unknown complex** data distribution p_{data}
- ▶ **Goal:** transform a **simple** data distribution p_0 in the **complex** p_{data}

Points $x_0 \sim p_0$ drawn from a **simple** data distribution $\mathcal{N}(0, 1)$

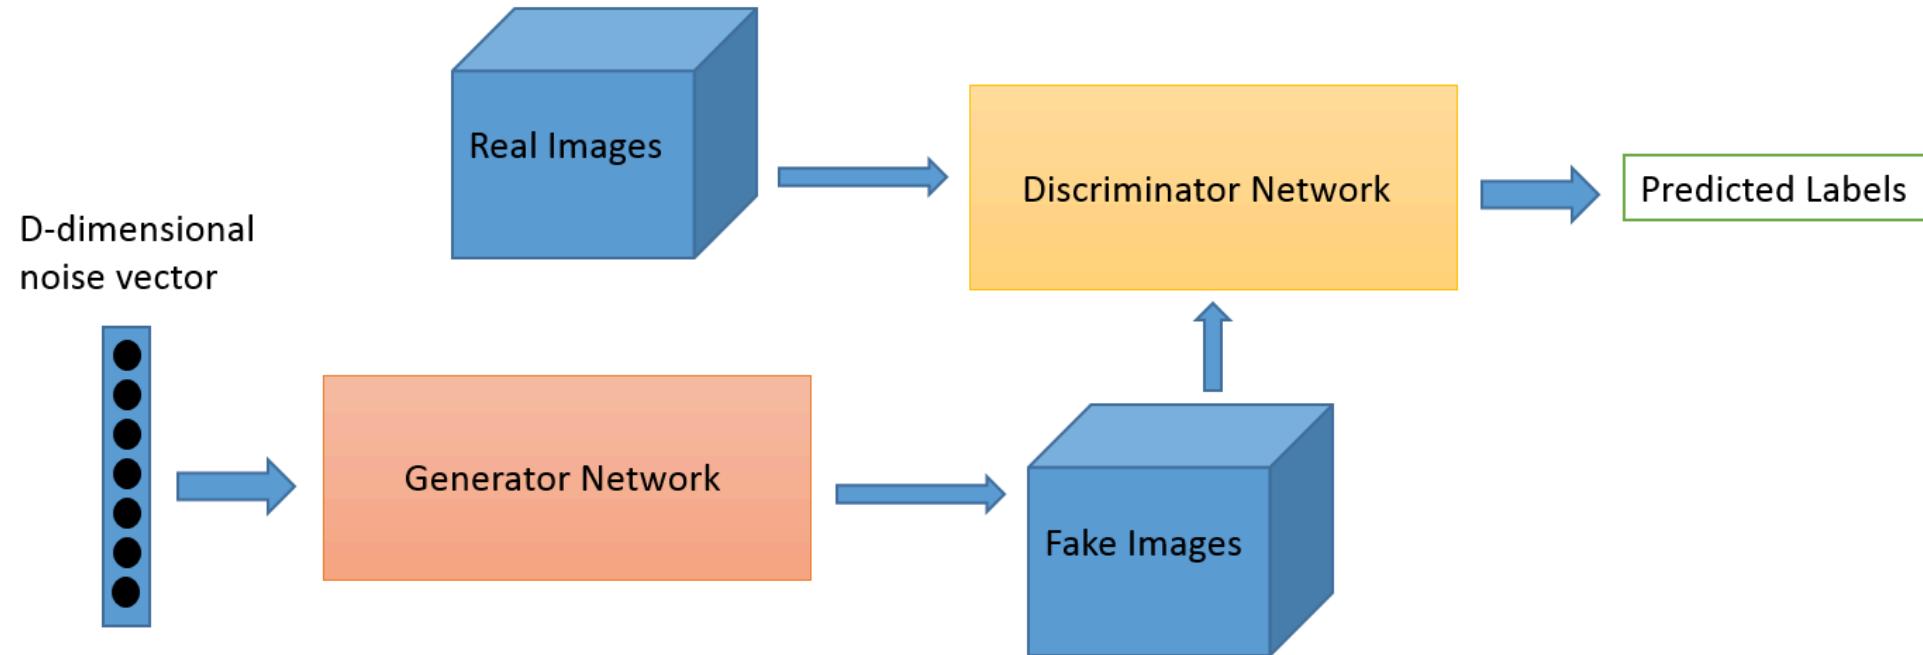


Complex data distribution p_{data}



Generative Adversarial Networks

Generative Adversarial Networks: General Idea



<https://datascience.eu/machine-learning/a-beginners-guide-to-generative-adversarial-networks-gans/>

Generative Adversarial Networks: Illustration

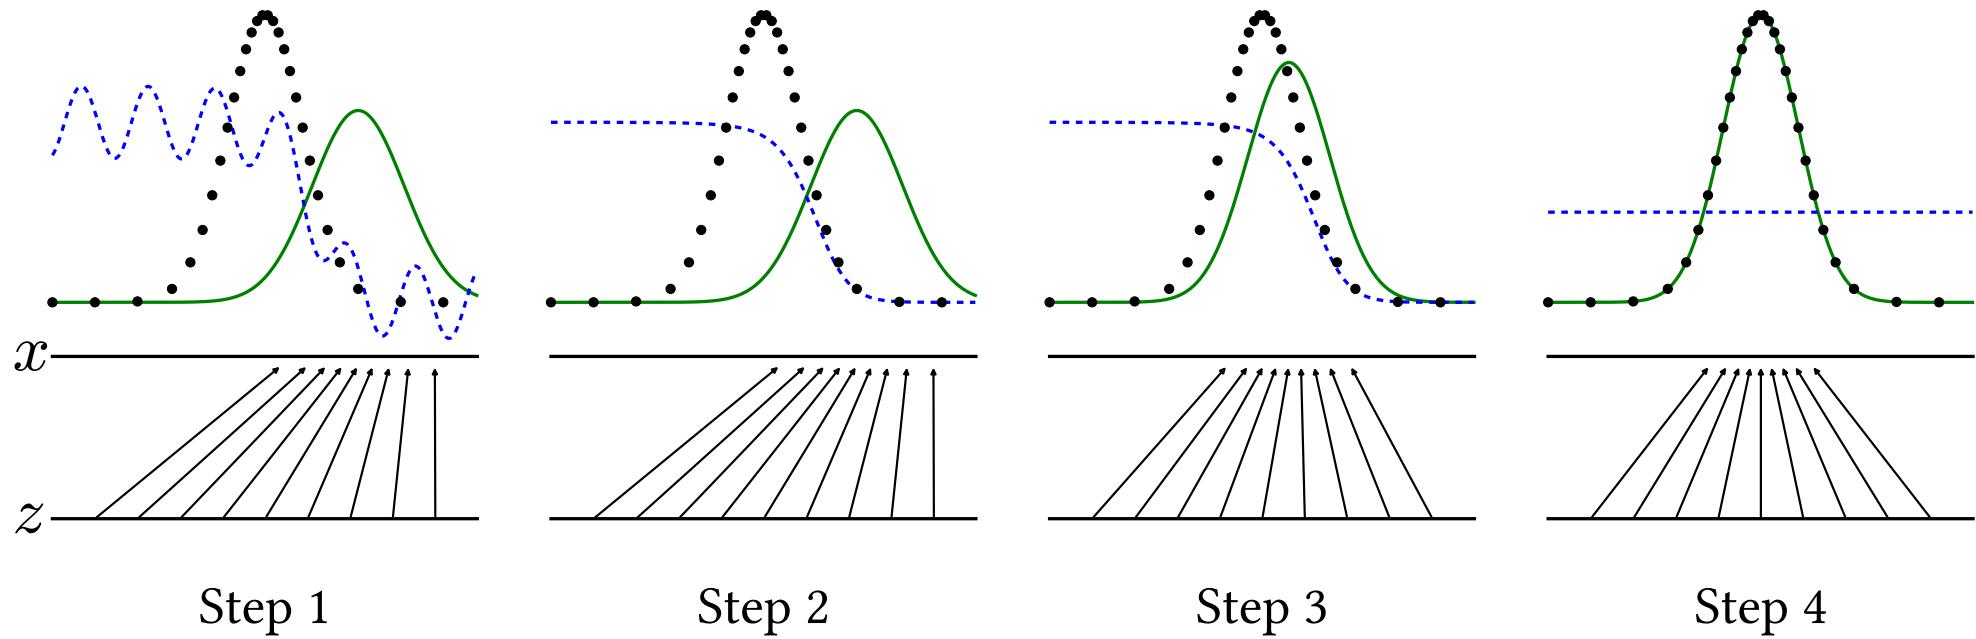


Image from [Goodfellow et al. \(2014\)](#)

How to Derive the GAN Loss?

Generative Adversarial Networks: Derive the Loss 1/2

We want to train:

Generative Adversarial Networks: Derive the Loss 1/2

We want to train:

- ▶ 2 networks

Generative Adversarial Networks: Derive the Loss 1/2

We want to train:

- ▶ 2 networks
 - A generator $G : \mathbb{R}^d \mapsto \mathbb{R}^d$

Generative Adversarial Networks: Derive the Loss 1/2

We want to train:

- ▶ 2 networks
 - A generator $G : \mathbb{R}^d \mapsto \mathbb{R}^d$
 - A discriminator $D : \mathbb{R}^d \mapsto [0, 1]$

Generative Adversarial Networks: Derive the Loss 1/2

We want to train:

- ▶ 2 networks
 - A generator $G : \mathbb{R}^d \mapsto \mathbb{R}^d$
 - A discriminator $D : \mathbb{R}^d \mapsto [0, 1]$
 - ▶ Tell how much an image is fake or not

Generative Adversarial Networks: Derive the Loss 1/2

We want to train:

- ▶ 2 networks
 - A generator $G : \mathbb{R}^d \mapsto \mathbb{R}^d$
 - A discriminator $D : \mathbb{R}^d \mapsto [0, 1]$
 - ▶ Tell how much an image is fake or not

Key Idea:

Generative Adversarial Networks: Derive the Loss 1/2

We want to train:

- ▶ 2 networks
 - A generator $G : \mathbb{R}^d \mapsto \mathbb{R}^d$
 - A discriminator $D : \mathbb{R}^d \mapsto [0, 1]$
 - ▶ Tell how much an image is fake or not

Key Idea:

- ▶ Maximum log-likelihood

Generative Adversarial Networks: Derive the Loss 1/2

We want to train:

- ▶ 2 networks
 - A generator $G : \mathbb{R}^d \mapsto \mathbb{R}^d$
 - A discriminator $D : \mathbb{R}^d \mapsto [0, 1]$
 - ▶ Tell how much an image is fake or not

Key Idea:

- ▶ Maximum log-likelihood
 - For a biased coin:

Generative Adversarial Networks: Derive the Loss 1/2

We want to train:

- ▶ 2 networks
 - A generator $G : \mathbb{R}^d \mapsto \mathbb{R}^d$
 - A discriminator $D : \mathbb{R}^d \mapsto [0, 1]$
 - ▶ Tell how much an image is fake or not

Key Idea:

- ▶ Maximum log-likelihood
 - For a biased coin:
 - ▶ $\mathbb{P}(\text{head}|D) = D$

Generative Adversarial Networks: Derive the Loss 1/2

We want to train:

- ▶ 2 networks
 - A generator $G : \mathbb{R}^d \mapsto \mathbb{R}^d$
 - A discriminator $D : \mathbb{R}^d \mapsto [0, 1]$
 - ▶ Tell how much an image is fake or not

Key Idea:

- ▶ Maximum log-likelihood
 - For a biased coin:
 - ▶ $\mathbb{P}(\text{head}|D) = D, \mathbb{P}(\text{tail}|D) = 1 - D$

Generative Adversarial Networks: Derive the Loss 1/2

We want to train:

- ▶ 2 networks
 - A generator $G : \mathbb{R}^d \mapsto \mathbb{R}^d$
 - A discriminator $D : \mathbb{R}^d \mapsto [0, 1]$
 - ▶ Tell how much an image is fake or not

Key Idea:

- ▶ Maximum log-likelihood
 - For a biased coin:
 - ▶ $\mathbb{P}(\text{head}|D) = D, \mathbb{P}(\text{tail}|D) = 1 - D$
 - ▶ Likelihood $\mathbb{P}(x|D) = D^x(1 - D)^{1-x}$

Generative Adversarial Networks: Derive the Loss 1/2

We want to train:

- ▶ 2 networks
 - A generator $G : \mathbb{R}^d \mapsto \mathbb{R}^d$
 - A discriminator $D : \mathbb{R}^d \mapsto [0, 1]$
 - ▶ Tell how much an image is fake or not

Key Idea:

- ▶ Maximum log-likelihood
 - For a biased coin:
 - ▶ $\mathbb{P}(\text{head}|D) = D, \mathbb{P}(\text{tail}|D) = 1 - D$
 - ▶ Likelihood $\mathbb{P}(x|D) = D^x(1 - D)^{1-x}$
 - ▶ Log-Likelihood $\log \mathbb{P}(x|D) = x \cdot \log D + (1 - x) \cdot \log(1 - D)$

Generative Adversarial Networks: Derive the Loss 1/2

We want to train:

- ▶ 2 networks
 - A generator $G : \mathbb{R}^d \mapsto \mathbb{R}^d$
 - A discriminator $D : \mathbb{R}^d \mapsto [0, 1]$
 - ▶ Tell how much an image is fake or not

Key Idea:

- ▶ Maximum log-likelihood
 - For a biased coin:
 - ▶ $\mathbb{P}(\text{head}|D) = D, \mathbb{P}(\text{tail}|D) = 1 - D$
 - ▶ Likelihood $\mathbb{P}(x|D) = D^x(1 - D)^{1-x}$
 - ▶ Log-Likelihood $\log \mathbb{P}(x|D) = x \cdot \log D + (1 - x) \cdot \log(1 - D)$

Maximum log-likelihood: for a given sequence of observations x_i

$$\max_D \sum_i x_i \cdot \log D + \sum_i (1 - x_i) \cdot \log(1 - D)$$

Generative Adversarial Networks: Derive the Loss 2/2

We want to train:

- ▶ 2 networks
 - A generator $G : \mathbb{R}^d \mapsto \mathbb{R}^d$
 - A discriminator $D : \mathbb{R}^d \mapsto [0, 1]$
 - ▶ Tell how much an image is fake or not

Generative Adversarial Networks: Derive the Loss 2/2

We want to train:

- ▶ 2 networks
 - A generator $G : \mathbb{R}^d \mapsto \mathbb{R}^d$
 - A discriminator $D : \mathbb{R}^d \mapsto [0, 1]$
 - ▶ Tell how much an image is fake or not

$$\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G}(\log(1 - D(\mathbf{x}_{\text{fake}})))$$

Generative Adversarial Networks: Derive the Loss 2/2

We want to train:

- ▶ 2 networks
 - A generator $G : \mathbb{R}^d \mapsto \mathbb{R}^d$
 - A discriminator $D : \mathbb{R}^d \mapsto [0, 1]$
 - ▶ Tell how much an image is fake or not

$$\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G}(\log(1 - D(\mathbf{x}_{\text{fake}})))$$

$$\max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G}(\log(1 - D(\mathbf{x}_{\text{fake}})))$$

Generative Adversarial Networks: Derive the Loss 2/2

We want to train:

- ▶ 2 networks
 - A generator $G : \mathbb{R}^d \mapsto \mathbb{R}^d$
 - A discriminator $D : \mathbb{R}^d \mapsto [0, 1]$
 - ▶ Tell how much an image is fake or not

$$\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} (\log(D(\mathbf{x}))) + \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G} (\log(1 - D(\mathbf{x}_{\text{fake}})))$$

$$\max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} (\log(D(\mathbf{x}))) + \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G} (\log(1 - D(\mathbf{x}_{\text{fake}})))$$

$$\max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} (\log(D(\mathbf{x}))) + \mathbb{E}_{x_{\text{noise}} \sim p_{\text{noise}}} (\log(1 - D(G(x_{\text{noise}}))))$$

Generative Adversarial Networks: Derive the Loss 2/2

We want to train:

- ▶ 2 networks
 - A generator $G : \mathbb{R}^d \mapsto \mathbb{R}^d$
 - A discriminator $D : \mathbb{R}^d \mapsto [0, 1]$
 - ▶ Tell how much an image is fake or not

$$\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G}(\log(1 - D(\mathbf{x}_{\text{fake}})))$$

$$\max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G}(\log(1 - D(\mathbf{x}_{\text{fake}})))$$

$$\max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{x_{\text{noise}} \sim p_{\text{noise}}}(\log(1 - D(G(x_{\text{noise}}))))$$

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{x_{\text{noise}} \sim p_{\text{noise}}}(\log(1 - D(G(x_{\text{noise}}))))$$

Generative Adversarial Networks: Derive the Loss 2/2

We want to train:

- ▶ 2 networks
 - A generator $G : \mathbb{R}^d \mapsto \mathbb{R}^d$
 - A discriminator $D : \mathbb{R}^d \mapsto [0, 1]$
 - ▶ Tell how much an image is fake or not

$$\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G}(\log(1 - D(\mathbf{x}_{\text{fake}})))$$

$$\max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G}(\log(1 - D(\mathbf{x}_{\text{fake}})))$$

$$\max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{x_{\text{noise}} \sim p_{\text{noise}}}(\log(1 - D(G(x_{\text{noise}}))))$$

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{x_{\text{noise}} \sim p_{\text{noise}}}(\log(1 - D(G(x_{\text{noise}}))))$$

What problem can you foresee?

Generative Adversarial Networks: Comment on the Loss

$$\min_G \max_D \mathcal{L}_{\text{GAN}}(G, D) := \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{x_{\text{noise}} \sim p_G}(\log(1 - D(G(x_{\text{noise}}))))$$

Generative Adversarial Networks: Comment on the Loss

$$\min_G \max_D \mathcal{L}_{\text{GAN}}(G, D) := \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{x_{\text{noise}} \sim p_G}(\log(1 - D(G(x_{\text{noise}}))))$$

Min-max problem: **very challenging**

Generative Adversarial Networks: Comment on the Loss

$$\min_G \max_D \mathcal{L}_{\text{GAN}}(G, D) := \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{x_{\text{noise}} \sim p_G}(\log(1 - D(G(x_{\text{noise}}))))$$

Min-max problem: **very challenging**

- ▶ Require specific algorithms

Generative Adversarial Networks: Comment on the Loss

$$\min_G \max_D \mathcal{L}_{\text{GAN}}(G, D) := \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{x_{\text{noise}} \sim p_G}(\log(1 - D(G(x_{\text{noise}}))))$$

Min-max problem: **very challenging**

- ▶ Require specific algorithms
- ▶ Tedious to tune

Generative Adversarial Networks: Comment on the Loss

$$\min_G \max_D \mathcal{L}_{\text{GAN}}(G, D) := \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{x_{\text{noise}} \sim p_G}(\log(1 - D(G(x_{\text{noise}}))))$$

Min-max problem: **very challenging**

- ▶ Require specific algorithms
- ▶ Tedious to tune

Comment:

- ▶ Closed-form formula for the max:

Generative Adversarial Networks: Comment on the Loss

$$\min_G \max_D \mathcal{L}_{\text{GAN}}(G, D) := \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{x_{\text{noise}} \sim p_G}(\log(1 - D(G(x_{\text{noise}}))))$$

Min-max problem: **very challenging**

- ▶ Require specific algorithms
- ▶ Tedious to tune

Comment:

- ▶ Closed-form formula for the max:

$$\arg \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{x_{\text{fake}} \sim p_G}(\log(1 - D(x_{\text{fake}}))) = \frac{p_{\text{data}}}{p_{\text{data}} + p_G}$$

Generative Adversarial Networks: Comment on the Loss

$$\min_G \max_D \mathcal{L}_{\text{GAN}}(G, D) := \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{x_{\text{noise}} \sim p_G}(\log(1 - D(G(x_{\text{noise}}))))$$

Min-max problem: **very challenging**

- ▶ Require specific algorithms
- ▶ Tedious to tune

Comment:

- ▶ Closed-form formula for the max:

$$\arg \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\log(D(\mathbf{x}))) + \mathbb{E}_{x_{\text{fake}} \sim p_G}(\log(1 - D(x_{\text{fake}}))) = \frac{p_{\text{data}}}{p_{\text{data}} + p_G}$$

Hence

$$\min_G \max_D \mathcal{L}_{\text{GAN}}(G, D) = \min_G \underbrace{\text{JS}}_{\text{compare distribution}}(p_{\text{data}} \parallel p_G)$$

Wasserstein Generative Adversarial Networks (WGANs): Idea

$$\min_G \max_D \mathcal{L}_{\text{GAN}}(G, D) = \min_G \underbrace{\text{JS}}_{\text{compares distribution}} (p_{\text{data}} \parallel p_G)$$

Wasserstein Generative Adversarial Networks (WGANs): Idea

$$\min_G \max_D \mathcal{L}_{\text{GAN}}(G, D) = \min_G \underbrace{\text{JS}}_{\text{compares distribution}} (p_{\text{data}} \parallel p_G)$$

WGANS Idea:

Wasserstein Generative Adversarial Networks (WGANs): Idea

$$\min_G \max_D \mathcal{L}_{\text{GAN}}(G, D) = \min_G \underbrace{\text{JS}}_{\text{compares distribution}}(p_{\text{data}} \parallel p_G)$$

WGANS Idea:

- ▶ replace JS, “Jensen Shanon” divergence

Wasserstein Generative Adversarial Networks (WGANs): Idea

$$\min_G \max_D \mathcal{L}_{\text{GAN}}(G, D) = \min_G \underbrace{\text{JS}}_{\text{compares distribution}}(p_{\text{data}} \parallel p_G)$$

WGANS Idea:

- ▶ replace JS, “Jensen Shanon” divergence
- ▶ With something else that compares distributions:

Wasserstein Generative Adversarial Networks (WGANs): Idea

$$\min_G \max_D \mathcal{L}_{\text{GAN}}(G, D) = \min_G \underbrace{\text{JS}}_{\text{compares distribution}} (p_{\text{data}} \parallel p_G)$$

WGANS Idea:

- ▶ replace JS, “Jensen Shanon” divergence
- ▶ With something else that compares distributions:
 - *e.g.*, Wasserstein Distance

Wasserstein Generative Adversarial Networks (WGANs): Idea

$$\min_G \max_D \mathcal{L}_{\text{GAN}}(G, D) = \min_G \underbrace{\text{JS}}_{\text{compares distribution}} (p_{\text{data}} \parallel p_G)$$

WGANS Idea:

- ▶ replace JS, “Jensen Shanon” divergence
- ▶ With something else that compares distributions:
 - *e.g.*, Wasserstein Distance
 - *e.g.*, Maximum Mean Discrepancy (MMD)

$$\min_G \max_D \mathcal{L}_{\text{WGAN}}(G, D) = \min_G \underbrace{\text{Wasserstein}}_{\text{compares distribution}} (p_{\text{data}}, p_G)$$

Wasserstein Generative Adversarial Networks (WGANs): Idea

$$\min_G \max_D \mathcal{L}_{\text{GAN}}(G, D) = \min_G \underbrace{\text{JS}}_{\text{compares distribution}} (p_{\text{data}} \parallel p_G)$$

WGANS Idea:

- ▶ replace JS, “Jensen Shanon” divergence
- ▶ With something else that compares distributions:
 - e.g., Wasserstein Distance
 - e.g., Maximum Mean Discrepancy (MMD)

$$\min_G \max_D \mathcal{L}_{\text{WGAN}}(G, D) = \min_G \underbrace{\text{Wasserstein}}_{\text{compares distribution}} (p_{\text{data}}, p_G)$$

Rewrite the Wasserstein distance, and after computations

$$\min_G \max_F \mathcal{L}_{\text{WGAN}}(G, F) = \min_G \max_{F, \text{Lip}(F) \leq 1} \mathbb{E}_{\textcolor{blue}{x} \sim p_{\text{data}}} F(\textcolor{blue}{x}) - \mathbb{E}_{\textcolor{red}{x}_{\text{fake}} \sim p_G} F(\textcolor{red}{x}_{\text{fake}})$$

GANs vs WGANs

$$\mathcal{L}_{\text{GAN}}(G, D) = \mathbb{E}_{\textcolor{blue}{x} \sim p_{\text{data}}} (\log(D(\textcolor{blue}{x}))) + \mathbb{E}_{\textcolor{red}{x}_{\text{fake}} \sim p_G} (\log(1 - D(\textcolor{red}{x}_{\text{fake}})))$$

GANs vs WGANs

$$\mathcal{L}_{\text{GAN}}(G, D) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} (\log(D(\mathbf{x}))) + \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G} (\log(1 - D(\mathbf{x}_{\text{fake}})))$$

$$\mathcal{L}_{\text{WGAN}}(G, F) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} F(\mathbf{x}) - \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G} F(\mathbf{x}_{\text{fake}})$$

GANs vs WGANs

$$\mathcal{L}_{\text{GAN}}(G, D) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} (\log(D(\mathbf{x}))) + \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G} (\log(1 - D(\mathbf{x}_{\text{fake}})))$$

$$\mathcal{L}_{\text{WGAN}}(G, F) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} F(\mathbf{x}) - \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G} F(\mathbf{x}_{\text{fake}})$$

- Additional “Lipschitz” constraint $\text{Lip}(F) \leq 1$

GANs vs WGANs

$$\mathcal{L}_{\text{GAN}}(G, D) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} (\log(D(\mathbf{x}))) + \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G} (\log(1 - D(\mathbf{x}_{\text{fake}})))$$

$$\mathcal{L}_{\text{WGAN}}(G, F) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} F(\mathbf{x}) - \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G} F(\mathbf{x}_{\text{fake}})$$

- Additional “Lipschitz” constraint $\text{Lip}(F) \leq 1$

How to enforce the constraint?

WGANS: How to Enforce the Constraint? $\text{Lip}(F) \leq 1/2$

Any Idea?

WGANS: How to Enforce the Constraint? $\text{Lip}(F) \leq 1/2$

Any Idea?

- ▶ Weight clipping
- ▶ Original WGANS paper ([Arjovsky, Chintala, and Bottou 2017](#))

WGANS: How to Enforce the Constraint? $\text{Lip}(F) \leq 1$ 1/2

Any Idea?

- ▶ Weight clipping
- ▶ Original WGANs paper ([Arjovsky, Chintala, and Bottou 2017](#))

```

1: while  $\theta$  has not converged do
2:   for  $t = 0, \dots, n_{\text{critic}}$  do
3:     Sample  $\{x^{(i)}\}_{i=1}^m \sim \mathbb{P}_r$  a batch from the real data.
4:     Sample  $\{z^{(i)}\}_{i=1}^m \sim p(z)$  a batch of prior samples.
5:      $g_w \leftarrow \nabla_w \left[ \frac{1}{m} \sum_{i=1}^m f_w(x^{(i)}) - \frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)})) \right]$ 
6:      $w \leftarrow w + \alpha \cdot \text{RMSProp}(w, g_w)$ 
7:      $w \leftarrow \text{clip}(w, -c, c)$ 
8:   end for
9:   Sample  $\{z^{(i)}\}_{i=1}^m \sim p(z)$  a batch of prior samples.
10:   $g_\theta \leftarrow -\nabla_\theta \frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)}))$ 
11:   $\theta \leftarrow \theta - \alpha \cdot \text{RMSProp}(\theta, g_\theta)$ 

```

WGANS: How to Enforce the Constraint? $\text{Lip}(F) \leq 1$ 2/2

$$\mathcal{L}_{\text{WGAN}}(G, F) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} F(\mathbf{x}) - \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G} F(\mathbf{x}_{\text{fake}})$$

WGANS: How to Enforce the Constraint? $\text{Lip}(F) \leq 1$ 2/2

$$\mathcal{L}_{\text{WGAN}}(G, F) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} F(\mathbf{x}) - \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G} F(\mathbf{x}_{\text{fake}})$$

- ▶ Additional “Lipschitz” constraint $\text{Lip}(F) \leq 1$

Idea to enforce the constraint?

WGANS: How to Enforce the Constraint? $\text{Lip}(F) \leq 1$ 2/2

$$\mathcal{L}_{\text{WGAN}}(G, F) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} F(\mathbf{x}) - \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G} F(\mathbf{x}_{\text{fake}})$$

- ▶ Additional “Lipschitz” constraint $\text{Lip}(F) \leq 1$

Idea to enforce the constraint?

- ▶ Gradient Penalty!
 - $\text{Lip}(F) \leq 1$ is equivalent to
 - $\| \nabla_x F \| \sim 1$

WGANS: How to Enforce the Constraint? $\text{Lip}(F) \leq 1$ 2/2

$$\mathcal{L}_{\text{WGAN}}(G, F) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} F(\mathbf{x}) - \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G} F(\mathbf{x}_{\text{fake}})$$

- ▶ Additional “Lipschitz” constraint $\text{Lip}(F) \leq 1$

Idea to enforce the constraint?

- ▶ Gradient Penalty!
 - $\text{Lip}(F) \leq 1$ is equivalent to
 - $\| \nabla_x F \| \sim 1$

$$\mathcal{L}_{\text{WGAN-gp}}(G, F) = \underbrace{\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} F(\mathbf{x}) - \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G} F(\mathbf{x}_{\text{fake}})}_{\mathcal{L}_{\text{WGAN}}(G, F)} + (\| \nabla_x F \| - 1)^2$$

WGANS: How to Enforce the Constraint? $\text{Lip}(F) \leq 1$ 2/2

$$\mathcal{L}_{\text{WGAN}}(G, F) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} F(\mathbf{x}) - \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G} F(\mathbf{x}_{\text{fake}})$$

- ▶ Additional “Lipschitz” constraint $\text{Lip}(F) \leq 1$

Idea to enforce the constraint?

- ▶ Gradient Penalty!
 - $\text{Lip}(F) \leq 1$ is equivalent to
 - $\| \nabla_x F \| \sim 1$

$$\mathcal{L}_{\text{WGAN-gp}}(G, F) = \underbrace{\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} F(\mathbf{x}) - \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G} F(\mathbf{x}_{\text{fake}})}_{\mathcal{L}_{\text{WGAN}}(G, F)} + (\| \nabla_x F \| - 1)^2$$

$$\mathcal{L}_{\text{WGAN-gp}}(G, F) = \underbrace{\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} F(\mathbf{x}) - \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G} F(\mathbf{x}_{\text{fake}})}_{\mathcal{L}_{\text{WGAN}}(G, F)} + \underbrace{\mathbb{E}_{x \sim p_{\text{data}} + p_G} (\| \nabla_x F(x) \| - 1)^2}_{\text{gradient penalty}}$$

LAB 1: WGANs

Pytorch recalls?

Pytorch 101 Before the LAB

```
def train_one_epoch(training_loader):
    for i, data in enumerate(training_loader):
        # Every data instance is an input + label pair
        inputs, labels = data

        # Zero your gradients for every batch!
        optimizer.zero_grad()

        # Make predictions for this batch
        outputs = model(inputs)

        # Compute the loss and its gradients
        loss = loss_fn(outputs, labels)
        loss.backward()

        # Adjust learning weights
        optimizer.step()
```

WGAN and WGAN-gp Losses

$$\min_G \max_F \mathcal{L}_{\text{WGAN}}(G, F) = \mathbb{E}_{\textcolor{blue}{x} \sim p_{\text{data}}} F(\textcolor{blue}{x}) - \mathbb{E}_{\textcolor{red}{x}_{\text{fake}} \sim p_G} F(\textcolor{red}{x}_{\text{fake}})$$

WGAN and WGAN-gp Losses

$$\min_G \max_F \mathcal{L}_{\text{WGAN}}(G, F) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} F(\mathbf{x}) - \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G} F(\mathbf{x}_{\text{fake}})$$

$$\mathcal{L}_{\text{WGAN-gp}}(G, F) = \underbrace{\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} F(\mathbf{x}) - \mathbb{E}_{\mathbf{x}_{\text{fake}} \sim p_G} F(\mathbf{x}_{\text{fake}})}_{\mathcal{L}_{\text{WGAN}}(G, F)} + \underbrace{\mathbb{E}_{x \sim p_{\text{data}} + p_G} (\| \nabla_x F(x) \| - 1)^2}_{\text{gradient penalty}}$$

References

- Arjovsky, Martin, Soumith Chintala, and Léon Bottou. 2017. “Wasserstein Generative Adversarial Networks”. In *International Conference on Machine Learning*, 214–23.
- Goodfellow, Ian, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. 2014. “Generative Adversarial Networks”. *Neurips*.
- Lam, Remi, Alvaro Sanchez-Gonzalez, Matthew Willson, Peter Wirsberger, Meire Fortunato, Ferran Alet, Suman Ravuri, et al. 2023. “Learning Skillful Medium-Range Global Weather Forecasting”. *Science* 382 (6677): 1416–21.
- Stability AI. 2023.
- Watson, Joseph L, David Juergens, Nathaniel R Bennett, Brian L Trippe, Jason Yim, Helen E Eisenach, Woody Ahern, et al. 2023. “De Novo Design of Protein Structure and Function with Rfdiffusion”. *Nature*, 1–3.

References

Xu, Minkai, Tomas Geffner, Karsten Kreis, Weili Nie, Yilun Xu, Jure Leskovec, Stefano Ermon, and Arash Vahdat. 2025. “Energy-Based Diffusion Language Models for Text Generation”. *ICLR*.